

```

import warnings
warnings.filterwarnings('ignore')
from sklearn.metrics import mean_squared_error, r2_score,
accuracy_score
from sklearn.linear_model import LinearRegression
import pandas as pd
import numpy as np
from plotnine import *
from sklearn.decomposition import PCA
from sklearn import metrics
from sklearn.linear_model import LogisticRegression # Logistic
Regression Model
from sklearn.preprocessing import StandardScaler #Z-score variables
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import accuracy_score, confusion_matrix,
plot_confusion_matrix
from sklearn.model_selection import train_test_split # simple TT split
cv
from sklearn.model_selection import KFold # k-fold cv
from sklearn.model_selection import LeaveOneOut #L00 cv
from sklearn.model_selection import cross_val_score # cross validation
metrics
from sklearn.model_selection import cross_val_predict # cross
validation metrics
from sklearn.cluster import DBSCAN
from sklearn.neighbors import NearestNeighbors
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import silhouette_score
rawData =
pd.read_csv("https://gist.githubusercontent.com/armgilles/194bcff35001
e7eb53a2a8b441e8b2c6/raw/92200bc0a673d5ce2110aaad4544ed6c4010f687/
pokemon.csv")
rawData.head()
newData = rawData.copy()
df = newData.copy()
df['Legendary'] = df['Legendary'].astype(int)
df

```

	#	Name	Type 1	Type 2	Total	HP	Attack
Defense \							
0	1	Bulbasaur	Grass	Poison	318	45	49
49							
1	2	Ivysaur	Grass	Poison	405	60	62
63							
2	3	Venusaur	Grass	Poison	525	80	82
83							
3	3	VenusaurMega Venusaur	Grass	Poison	625	80	100
123							

4	4	Charmander	Fire	NaN	309	39	52
43							
..
...							
795	719	Diancie	Rock	Fairy	600	50	100
150							
796	719	DiancieMega Diancie	Rock	Fairy	700	50	160
110							
797	720	HoopaHoopa Confined	Psychic	Ghost	600	80	110
60							
798	720	HoopaHoopa Unbound	Psychic	Dark	680	80	160
60							
799	721	Volcanion	Fire	Water	600	80	110
120							

	Sp. Atk	Sp. Def	Speed	Generation	Legendary
0	65	65	45	1	0
1	80	80	60	1	0
2	100	100	80	1	0
3	122	120	80	1	0
4	60	50	65	1	0
..
795	100	150	50	6	1
796	160	110	110	6	1
797	150	130	70	6	1
798	170	130	80	6	1
799	130	90	70	6	1

[800 rows x 13 columns]

1)

When creating a Logistic Regression model is it best to use HP, Attack, Defense, and Speed or Special Attack and Special Defense as predictors to determine if a Pokemon is Legendary?

The accuracy scores of the models were 90.63% and 91.88% showing that both models were very effective. However in the future I can create a better model for HP, Attack, Defense, and Speed if I don't use Attack as a predictor because when graphing the coefficients you can see that Attack has the lowest coefficient of 0.46 showing us that it has the lowest impact on determining if a Pokemon is Legendary.

However, I would definitely go with my second model to make prediction because with only 2 predictors it was just as accurate as the first model.

This model differed from my analysis to determine which predictor groups were the most accurate rather than which specific predictors.


```
predictors = ["HP", "Attack", "Defense", "Speed"]
X_train, X_test, y_train, y_test = train_test_split(df[predictors],
```

```

df["Legendary", test_size=0.2)
z = StandardScaler()
z.fit(X_train[predictors])
X_test[predictors]=z.transform(X_test[predictors])
X_train[predictors]=z.transform(X_train[predictors])
lr = LogisticRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_train)
confusion_matrix(y_train, y_pred)
X = df[predictors]
y = df["Legendary"]
plot_confusion_matrix(lr,X_test,y_test)

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7f5d372d6190>

```



```

print("Accuracy score:",accuracy_score(y_test, lr.predict(X_test)))

```

Accuracy score: 0.90625

```

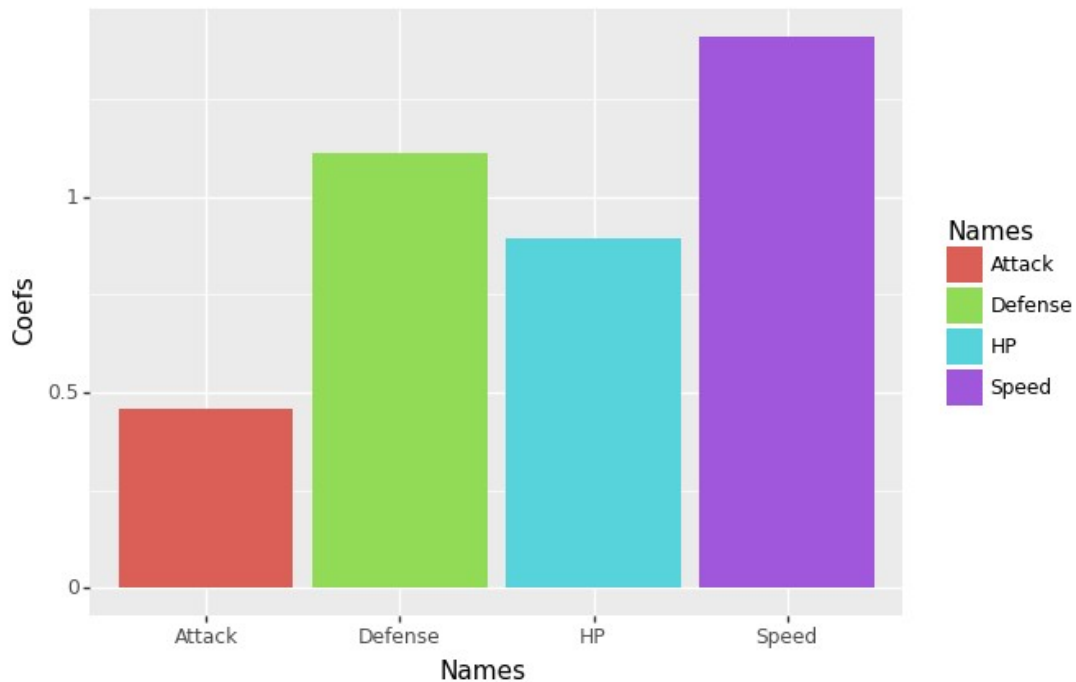
coef = pd.DataFrame({"Coefs": lr.coef_[0],
                     "Names": predictors})

```

coef

	Coefs	Names
0	0.893368	HP
1	0.459688	Attack
2	1.112682	Defense
3	1.408566	Speed


```
(ggplot(coef, aes(x = "Names", y = "Coefs", fill = "Names" )) +
  geom_bar(stat = "identity"))
```



```
<ggplot: (8752395868793)>
```

```
predictors2 = ["Sp. Atk","Sp. Def"]
X_train, X_test, y_train, y_test = train_test_split(df[predictors2],
df["Legendary"], test_size=0.2)
z = StandardScaler()
z.fit(X_train[predictors2])
X_test[predictors2]=z.transform(X_test[predictors2])
X_train[predictors2]=z.transform(X_train[predictors2])
lr = LogisticRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_train)
confusion_matrix(y_train, y_pred)
X = df[predictors2]
y = df["Legendary"]
plot_confusion_matrix(lr,X_test,y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7f5d37103ad0>
```



```
print("Accuracy score:",accuracy_score(y_test, lr.predict(X_test)))
```


Accuracy score: 0.91875

```
coef = pd.DataFrame({"Coefs": lr.coef_[0],
                    "Names": predictors2})
```

```
coef
```

	Coefs	Names
0	1.365821	Sp. Atk
1	0.861780	Sp. Def

```
(ggplot(coef, aes(x = "Names", y = "Coefs", fill = "Names" )) +
  geom_bar(stat = "identity"))
```



<ggplot: (8752395855841)>

2)

When using Total and Sp.Attack (the most influential predictor for Legendaries) to cluster Legendaries what is the best clustering model to use?

When creating clustering models I determined that the best model to use was KNN because KNN ended up having an accuracy score of 93.13% which could be because my model was very accurate in counting how many neighbors the different Pokemon had to classify them.

DBSCAN was very inaccurate because the model was very random and did not have too much noise, so I was not surprised to see that the silhouette score was a low 0.46, showing that the model clustered badly.

Gaussian Mixture method was the most inefficient, possibly because the data was so random, and had a lower silhouette score of 0.175.

Here I used DBSCAN to cluster Legendaries, and ended up realizing that the randomness of the data and lack of noise points were really bad for my model.

In my analysis I wantd to use DBSCAN to cluster types, but I decided to cluster legendaries for accuracy, and when DBSCAN didn't work I tried other methods.

```
mins = 4
predictors = ["HP", "Attack", "Defense", "Speed", "Sp. Atk", "Sp. Def", "Total"]
nn = NearestNeighbors(n_neighbors = mins + 1)
X = df.copy()
```

```

X = X[predictors]
z = StandardScaler()
X["predictors"] = z.fit(X[predictors])
nn.fit(X[predictors])
distances, neighbors = nn.kneighbors(X[predictors])

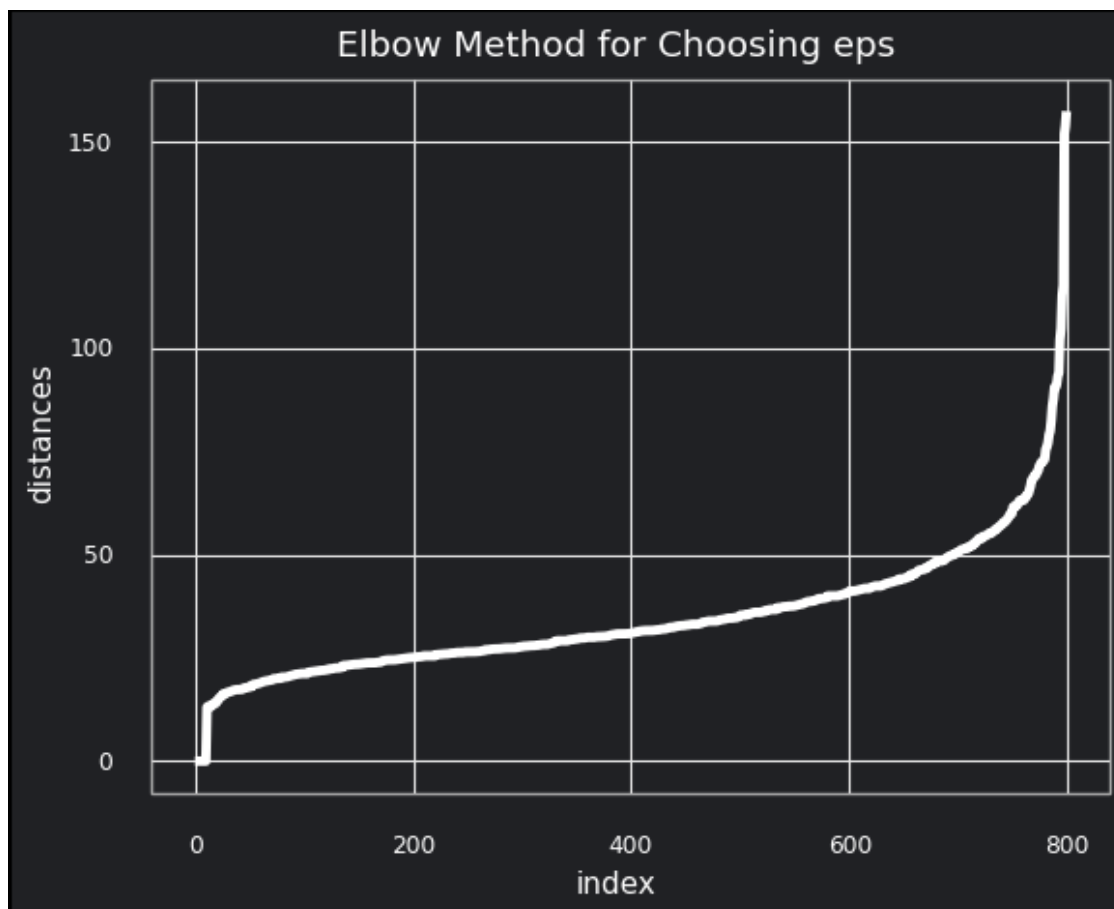
distances

distances = np.sort(distances[:, mins], axis = 0)

distances_pokemon = pd.DataFrame({"distances": distances,
                                   "index": list(range(0, len(distances)))})
plt = (ggplot(distances_pokemon, aes(x = "index", y = "distances")) +
       geom_line(color = "white", size = 2) + theme_minimal() +
       labs(title = "Elbow Method for Choosing eps") +
       theme(panel_grid_minor = element_blank(),
             rect = element_rect(fill = "#202124ff"),
             axis_text = element_text(color = "white"),
             axis_title = element_text(color = "white"),
             plot_title = element_text(color = "white"),
             panel_border = element_line(color = "darkgray"),
             plot_background = element_rect(fill = "#202124ff")
       ))
ggsave(plot=plt, filename='elbow.png', dpi=300)

plt

```




```
<ggplot: (8752395868925)>
```

```
df1 = DBSCAN(eps = 50, min_samples = 5).fit(X[predictors])
labsList = ["Noise"]
labsList = labsList + ["Cluster " + str(i) for i in
range(1,len(set(df1.labels_)))]

X["assignments"] = df1.labels_

(ggplot(X, aes(x = "Total", y = "Sp. Atk", color =
"factor(assignments)"))
+ geom_point()
+ theme_minimal()
+ scale_color_discrete(name = "Cluster Assignment", labels = labsList)

+ theme(panel_grid_major = element_blank())
+ labs(title = "DBSCAN: eps = 55, min_samples = 5"))
```

```
<ggplot: (8752400466557)>
```

```
ss_dbscan = X.loc[(X.assignments >= 0)]
print("SILHOUETTE:", silhouette_score(ss_dbscan[["Total", "Sp. Atk"]],
ss_dbscan["assignments"]))
```

```
SILHOUETTE: 0.4607780894566231
```

Here I used a Gaussian Mixture model to determine the probability that each Pokemon is a Legendary, and I determined that it was ineffective because of how random the model was.

```
EM = GaussianMixture(n_components = 2)
X = df.copy()
X = X[predictors]
z = StandardScaler()
```


```
X[predictors] = z.fit_transform(X)
EM.fit(X)
```

```
cluster = EM.predict(X)
```

```
X["cluster"] = cluster
```

```
print((ggplot(X, aes(x = "Total", y = "Sp. Atk", color = "cluster")) +
geom_point()))
```

```
print("SILHOUETTE: ", silhouette_score(X, cluster))
```



```
<ggplot: (8752396197761)>
SILHOUETTE: 0.17481438594257187
```

Here I used KNN and determined that it was super efficient because it tracked how many neighbors each Pokemon had.

```
predictors = ["Total", "Sp. Atk"]
X = df[predictors]
y = df["Legendary"]
```

```
n_neighbors = 5
```

```
knn = KNeighborsClassifier(n_neighbors = n_neighbors)
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size =
0.2)
```

```
z = StandardScaler()
z.fit(X_train)
X_train[predictors] = z.transform(X_train)
X_test[predictors] = z.transform(X_test)
```

```
knn.fit(X_train,y_train)
```

```
X_train
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
```

```

def plotKNN2D(Xdf,y,k):
    # X can only have 2 dimensions becuae of plotting
    print("-----")
    print(Xdf.columns)

    Xdf.columns = ["x0", "x1"]

    #grab the range of features for each feature
    x0_range = np.linspace(min(Xdf["x0"]) - np.std(Xdf["x0"]),
                           max(Xdf["x0"]) + np.std(Xdf["x0"]), num =
100)
    x1_range = np.linspace(min(Xdf["x1"]) - np.std(Xdf["x1"]),
                           max(Xdf["x1"]) + np.std(Xdf["x1"]), num =
100)

    #get all possible points on graph
    x0 = np.repeat(x0_range,100)
    x1 = np.tile(x1_range,100)
    x_grid = pd.DataFrame({"x0": x0, "x1": x1})

    #build model
    knn = KNeighborsClassifier(n_neighbors = k)
    knn.fit(Xdf,y)


    # bredict all background points
    p = knn.predict(x_grid)
    x_grid["p"] = p #add to dataframe

    #build the plot
    bound = (ggplot(x_grid, aes(x = "x0", y = "x1", color =
"factor(p)")) +
              geom_point(alpha = 0.2, size = 0.2) + theme_minimal()
+
              scale_color_discrete(name = "Class") +
              geom_point(data = Xdf, mapping = aes(x = "x0", y =
"x1", color = "factor(y)", size = 2))
    print(bound)


plotKNN2D(X_train,y_train, k = n_neighbors)

-----
Index(['Total', 'Sp. Atk'], dtype='object')

```



```
<ggplot: (8752403684845)>
print("ACCURACY:", knn.score(X_test, y_test))
ACCURACY: 0.93125
plot_confusion_matrix(knn, X_test, y_test)
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7f5d3b99bd90>
```



3)

When using HP, Attack, Defense, Speed, Special attack, Special Defense, and Total as my predictors to determine what type a Pokemon is, can I use PCA to create a more efficient model?

In this model I determined that when using PCA to remove 4 of my predictors, and using a regular Logistic Regression both models produced very inaccurate predictions.

In my PCA model I determined that I can keep over 90% accuracy by keeping only 4 predictors and produced a model with an R2 score of -1 for train and -1.47 for test and an MSE of 59 for train and 71 for test, showing that the variance of the data could not be explained by the model, and that the errors squared were, and that the regression line was nowhere near my points.

In my Logistic Regression model I was able to determine that the R2 score was -0.89 for train and -1.4 for test and the MSE was 56 for train and 69.67 for test, so this was also very inaccurate. When graphing my coefficients I determined that none of the coefficients were above an absolute value of 0.1, meaning that none of the predictors were impactful at all.

```
features = ["HP", "Attack", "Defense", "Speed", "Sp. Atk", "Sp. Def", "Total"]
features
cleanup_nums = {"Type 1": {"Water": 0, "Steel": 1, "Rock": 2, "Psychic": 3, "Poison": 4, "Normal": 5, "Ice": 6, "Ground": 7, "Grass": 8, "Ghost": 9, "Flying": 10, "Fire": 11, "Fighting": 12, "Fairy": 13, "Electric": 14, "Dragon": 15, "Dark": 16, "Bug": 16}}
df = df.replace(cleanup_nums)
X_train, X_test, y_train, y_test = train_test_split(df[features],
```

```

df["Type 1"], test_size=0.2)
z = StandardScaler()
X_train[predictors] = z.fit_transform(X_train[predictors])
X_test[predictors] = z.transform(X_test[predictors])

lr = LogisticRegression()

lr.fit(X_train, y_train)


pca = PCA()
pca.fit(X_train[features], y_train)
pcaDF = pd.DataFrame({"expl_var" :
                      pca.explained_variance_ratio_,
                      "pc": range(1,8),
                      "cum_var":
                      pca.explained_variance_ratio_.cumsum()
                      })

pcaDF.head()

  expl_var  pc  cum_var
0  0.482405   1  0.482405
1  0.197129   2  0.679534
2  0.145420   3  0.824954
3  0.117848   4  0.942801
4  0.057047   5  0.999848


(ggplot(pcaDF, aes(x = "pc", y = "expl_var")) + geom_line() +
geom_point())

```



```
<ggplot: (8752400579529)>
```

```
(ggplot(pcaDF, aes(x = "pc", y = "cum_var")) + geom_line(color =  
"pink") +  
geom_point(color = "pink") + geom_hline(yintercept = 0.90))
```



```
<ggplot: (8752400564509)>
```

Here I used PCA to create my model with only 4 of the predictors.

```
pcomps4Train = pca.transform(X_train)
pcomps4Test = pca.transform(X_test)
pcomps4Train = pd.DataFrame(pcomps4Train[:, 0:4])
pcomps4Test = pd.DataFrame(pcomps4Test[:, 0:4])
lr1 = LogisticRegression()
lr1.fit(pcomps4Train, y_train)
```

```
print("MSE Train:", mean_squared_error(y_train,
lr1.predict(pcomps4Train)))
print("MSE Test:", mean_squared_error(y_test,
lr1.predict(pcomps4Test)))
```

```
print("R2 Train:", r2_score(y_train, lr1.predict(pcomps4Train)))
print("R2 Test:", r2_score(y_test, lr1.predict(pcomps4Test)))
```

```
MSE Train: 59.4140625
MSE Test: 71.28125
R2 Train: -1.0009078741266477
R2 Test: -1.4690957112857492
```


Here I used all my predictors for a Logistic Regression.

```
yPredTrain = lr.predict(X_train)
yPredTest = lr.predict(X_test)
```

```
print("MSE Train:" , mean_squared_error(y_train,yPredTrain ))
print("MSE Test:" , mean_squared_error(y_test,yPredTest ))
```

```
print("R2 Train:",r2_score(y_train,yPredTrain ))
print("R2 Test:",r2_score(y_test,yPredTest ))
```

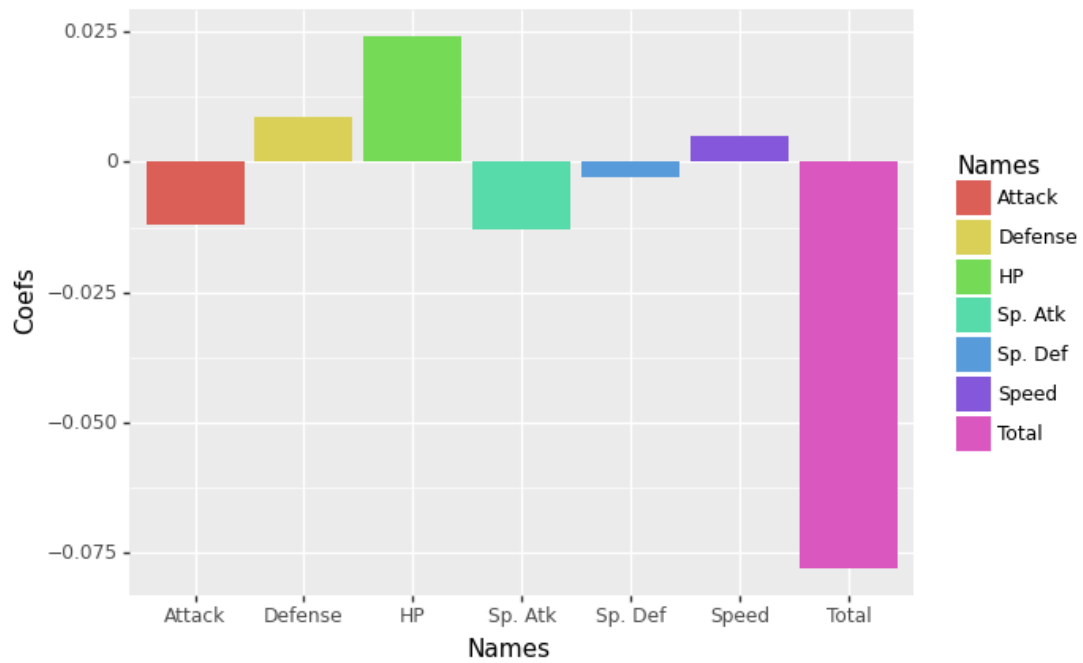
```
MSE Train: 56.1140625
MSE Test: 69.66875
R2 Train: -0.8897726359897515
R2 Test: -1.413240674590289
```

```
coef = pd.DataFrame({"Coefs": lr.coef_[0],
                     "Names": features})
```

```
coef
```

	Coefs	Names
0	0.024047	HP
1	-0.011997	Attack
2	0.008517	Defense
3	0.004989	Speed
4	-0.013052	Sp. Atk
5	-0.002963	Sp. Def
6	-0.078013	Total

```
(ggplot(coef, aes(x = "Names", y = "Coefs", fill = "Names" )) +
geom_bar(stat = "identity"))
```



<ggplot: (8752400496581)>